

Prediction of Life Expectancy using Regression

David

2024-09-19

The main objective of this analysis is to build a Regression model that can accurately predict Life Expectancy for any country in Africa. I'm mainly focusing on Africa because Africa has the highest number of least developed countries. These are low-income countries which are highly vulnerable to economic shocks and have low levels of human assets (UN, 2023). These countries are mainly characterized by slow economic growth and high population growth, which has led to the rising number of people living in extreme poverty.

```
# Load packages
suppressMessages(
  {library(tidyverse)
   library(caret)
   library(janitor)
   library(mlr)
   library(psych)
   library(parallel)
   library(parallelMap)
  }
)

# Import data
Life_Expectancy <- read.csv("Life_Expectancy.csv")

# View the structure of the dataset
Life_Expectancy |> glimpse()

## Rows: 1,904
## Columns: 17
## $ Country                <chr> "Albania",
## $ Year                    <int> 2000, 2000,
## $ Continent               <chr> "Europe",
## $ Least.Developed         <lgl> FALSE, FALSE,
## $ Life.Expectancy         <dbl> 73.95500,
## $ Population              <int> 3089027,
## $ CO2.emissions           <dbl> 1.02621311,
## $ ...                     <dbl> 2.5787...
```

```
## $ Health.expenditure <dbl> 7.233370,
3.489033...
## $ Electric.power.consumption <dbl> 1414.70378,
652.80...
## $ Forest.area <dbl> 28.0766423,
0.6629...
## $ GDP.per.capita <dbl> 3860.8046,
8436.99...
## $ Individuals.using.the.Internet <dbl> 0.114097347,
0.491...
## $ Military.expenditure <dbl> 1.2463602,
3.43338...
## $ People.practicing.open.defecation <dbl> 0.88885278,
6.2339...
## $ People.using.at.least.basic.drinking.water.services <dbl> 86.75447,
89.84883...
## $ Obesity.among.adults <dbl> 12.8, 15.0,
3.0, 2...
## $ Beer.consumption.per.capita <dbl> 1.33431,
0.11978, ...
```

The data has 1904 observations of 17 variables. Country and Continent are character variables, least developed is a logical variable while the rest of the variables are numeric.

Clean the data

```
# Clean variable names
Life_Expectancy <- clean_names(Life_Expectancy)

# Rename variables with longer variable names to have shorter names
Life_Expectancy <- Life_Expectancy |> rename(open_defecation =
"people_practicing_open_defecation",
      basic_drinking_water =
"people_using_at_least_basic_drinking_water_services",
      adults_obesity = "obesity_among_adults",
      beer_consumption = "beer_consumption_per_capita",
      internet = "individuals_using_the_internet")

# Check for missing values
map_dbl(Life_Expectancy, ~length(which(is.na(.))))

##           country           year
##           0             0
##      continent least_developed
##           0             0
##    life_expectancy      population
##           0             0
##      co2_emissions health_expenditure
##           0             0
## electric_power_consumption      forest_area
```

```
##           0           0
##      gdp_per_capita      internet
##           0           0
##      military_expenditure  open_defecation
##           0           0
##      basic_drinking_water  adults_obesity
##           0           0
##      beer_consumption
##           0
```

There are no missing values in the data.

```
# Check for duplicated observations
sum(duplicated(Life_Expectancy))

## [1] 0
```

There are no duplicated observations in the data as well.

Exploratory Data Analysis

I'm more interested in Africa, so I will filter the data for Africa. I'd like to obtain the average Life Expectancy in Africa between the period 2000 to 2015, average GDP per Capita, Health Expenditure, the population with access to clean water services and the population which is still practicing open defecation.

```
# Filter the data for Africa
Africa_Data <- Life_Expectancy |> filter( continent == "Africa")
```

The data for Africa has 448 observations.

```
# Calculate summary statistics
Africa_Data |> select(-c(country, continent)) |> summary()

##      year      least_developed life_expectancy  population
##  Min.   :2000   Mode :logical   Min.   :43.06   Min.   : 1186873
## 1st Qu.:2004   FALSE:256      1st Qu.:53.49   1st Qu.: 6037752
## Median :2008   TRUE :192      Median :58.43   Median : 18429400
## Mean   :2008                      Mean   :59.72   Mean   : 28169437
## 3rd Qu.:2011                      3rd Qu.:64.22   3rd Qu.: 36229780
## Max.   :2015                      Max.   :76.09   Max.   :181137454
## co2_emissions  health_expenditure electric_power_consumption
##  Min.   :0.03242   Min.   : 0.410   Min.   : 22.76
## 1st Qu.:0.25974   1st Qu.: 3.561   1st Qu.: 102.25
## Median :0.57964   Median : 4.466   Median : 215.42
## Mean   :1.54075   Mean   : 4.796   Mean   : 727.14
## 3rd Qu.:2.17482   3rd Qu.: 5.672   3rd Qu.:1019.67
## Max.   :9.60845   Max.   :10.716   Max.   :4851.69
## forest_area    gdp_per_capita      internet
## military_expenditure
```

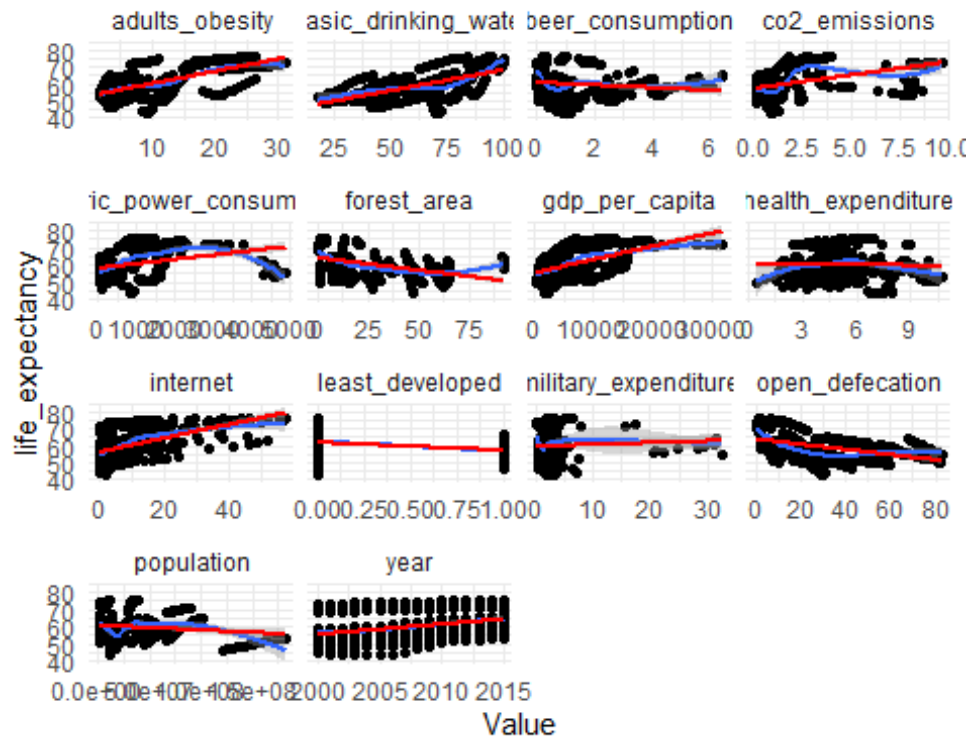
## Min. : 0.0521	Min. : 434.8	Min. : 0.00101	Min. : 0.1427
## 1st Qu.: 9.0794	1st Qu.: 1828.3	1st Qu.: 0.96989	1st Qu.: 1.0555
## Median :21.5803	Median : 3428.9	Median : 3.97500	Median : 1.4898
## Mean :28.1303	Mean : 5826.6	Mean : 8.60403	Mean : 2.8765
## 3rd Qu.:45.7896	3rd Qu.: 8804.6	3rd Qu.:11.52500	3rd Qu.: 2.7355
## Max. :91.9781	Max. :31702.9	Max. :57.08000	Max. :32.6557
## open_defecation	basic_drinking_water	adults_obesity	beer_consumption
## Min. : 0.08106	Min. :18.09	Min. : 1.600	Min. :0.0000
## 1st Qu.: 7.47031	1st Qu.:49.20	1st Qu.: 4.300	1st Qu.:0.2396
## Median :19.08088	Median :64.81	Median : 6.400	Median :0.6300
## Mean :25.66821	Mean :65.63	Mean : 9.582	Mean :1.1585
## 3rd Qu.:41.58956	3rd Qu.:82.50	3rd Qu.:12.550	3rd Qu.:1.4234
## Max. :82.72894	Max. :99.87	Max. :31.000	Max. :6.4100

- The average Life Expectancy in Africa between 2000 to 2015 was 59.79 years (SD = 8.12). The average population was 28,169,437.
- The average GDP Per Capita for Africa (the average income per person) between 2000 to 2015 was about \$5,826.58, the average Health Expenditure was 4.796% of GDP, while the average military expenditure was 2.88% of GDP.
- On average, 65.63% of the population have access to basic drinking water services, but 25.67% of the population still practice open defecation. This is still a major challenge for Africa.
- Also on average, 8.6% of the population have access to internet connection, and 9.58% of the population are obese.

```
# Convert the data to Long format for plotting
UntidyData <- gather(Africa_Data, key = "Variable", value = "Value",
                     -c(country, continent, life_expectancy))

# Plot
ggplot(UntidyData, aes(Value, life_expectancy)) +
  facet_wrap(~Variable, scale = "free_x")+
  geom_point() +
  geom_smooth() +
  geom_smooth(method = "lm", col = "red") +
  theme_minimal()

## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'
```



Most of the predictor variables have non-linear relationship with the outcome variable. Health expenditure and military expenditure don't seem to have a relationship with Life Expectancy.

Model training

```
# Select the variables to be used for model training (omit country &
continent)
data <- Africa_Data |> select(-c(country, continent))

## Partition the data into training and test sets

# Set seed for reproducibility
set.seed(42)
# Partition the data
splitIndex <- createDataPartition(data$life_expectancy, p = 0.80, list =
FALSE)
# Assign 80% to training set
training_data <- data[splitIndex, ]
# Assign test set the remaining 20%
test_data <- data[-splitIndex, ]

# Convert all the features in training data to numeric
training_data <- mutate_all(training_data, .funs = ~ as.numeric(.))
# Convert all the features in test data to numeric
test_data <- mutate_all(test_data, .funs = ~ as.numeric(.))
```

Linear Regression model

```
# Define regression task
reg_task <- makeRegrTask(data = training_data, target = "life_expectancy")
# Define Learner
reg_learner <- makeLearner("regr.lm")


# Estimate variable importance using linear correlation
filterVals <- generateFilterValuesData( reg_task, method =
"linear.correlation")

# Get the correlation coefficients
filterVals$data
```

	name	type	filter	value
	<char>	<char>	<fctr>	<num>
## 1:	year	numeric	linear.correlation	0.30431729
## 2:	least_developed	numeric	linear.correlation	0.25650947
## 3:	population	numeric	linear.correlation	0.10987719
## 4:	co2_emissions	numeric	linear.correlation	0.41608276
## 5:	health_expenditure	numeric	linear.correlation	0.02896530
## 6:	electric_power_consumption	numeric	linear.correlation	0.32426011
## 7:	forest_area	numeric	linear.correlation	0.42050964
## 8:	gdp_per_capita	numeric	linear.correlation	0.55367524
## 9:	internet	numeric	linear.correlation	0.55811083
## 10:	military_expenditure	numeric	linear.correlation	0.07339473
## 11:	open_defecation	numeric	linear.correlation	0.41195536
## 12:	basic_drinking_water	numeric	linear.correlation	0.61581017
## 13:	adults_obesity	numeric	linear.correlation	0.65227749
## 14:	beer_consumption	numeric	linear.correlation	0.12078963

Most of the predictors are moderately correlated with the outcome variable. Population, Health Expenditure, Military Expenditure and beer consumption are weakly correlated with life expectancy. None of the predictors is highly correlated with the outcome variable.

```
# Plot the feature importance
plotFilterValues(filterVals) + theme_bw() + coord_flip()
```



Despite Health Expenditure being a very important determinant of life expectancy, it contributes very little information to the outcome variable. This means that Health is underfunded in Africa and this remains a major concern. Adults obesity, basic drinking water services, access to internet and GDP per Capita, respectively, contribute the highest information to the outcome variable. Open defecation, CO2 emissions and forest area cover contribute nearly the same amount of information to the outcome variable.

From my EDA, most of the predictor variables didn't have a linear relationship with the outcome Life Expectancy, but I'll use all the variables for modelling because they might be of importance to non-linear models.

```
# Fit the Linear Regression model
Linear_model <- train(reg_learner, reg_task)
# Obtain model results
results <- getLearnerModel(Linear_model)

# Have a model summary
summary(results)

##
## Call:
## stats::lm(formula = f, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.7832 -2.2830  0.1162  2.0636  9.1261
##
```

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.367e+02  1.124e+02  -1.216   0.2247
## year        9.699e-02  5.624e-02   1.725   0.0855 .
## least_developed 6.990e+00  7.478e-01   9.348 < 2e-16 ***
## population    -3.582e-08  7.321e-09  -4.892 1.53e-06 ***
## co2_emissions -2.153e-01  3.434e-01  -0.627   0.5311
## health_expenditure -6.609e-01  1.565e-01  -4.222 3.10e-05 ***
## electric_power_consumption -2.683e-03  5.813e-04  -4.615 5.55e-06 ***
## forest_area    -1.735e-01  1.558e-02 -11.134 < 2e-16 ***
## gdp_per_capita  3.973e-04  8.927e-05   4.451 1.16e-05 ***
## internet       1.642e-01  2.735e-02   6.005 4.86e-09 ***
## military_expenditure 1.488e-01  3.957e-02   3.759  0.0002 ***
## open_defecation -1.440e-01  1.516e-02  -9.496 < 2e-16 ***
## basic_drinking_water 1.022e-01  2.221e-02   4.599 5.96e-06 ***
## adults_obesity  3.137e-01  6.339e-02   4.949 1.17e-06 ***
## beer_consumption -6.205e-02  2.542e-01  -0.244   0.8073
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.55 on 345 degrees of freedom
## Multiple R-squared:  0.8142, Adjusted R-squared:  0.8067
## F-statistic: 108 on 14 and 345 DF, p-value: < 2.2e-16
```

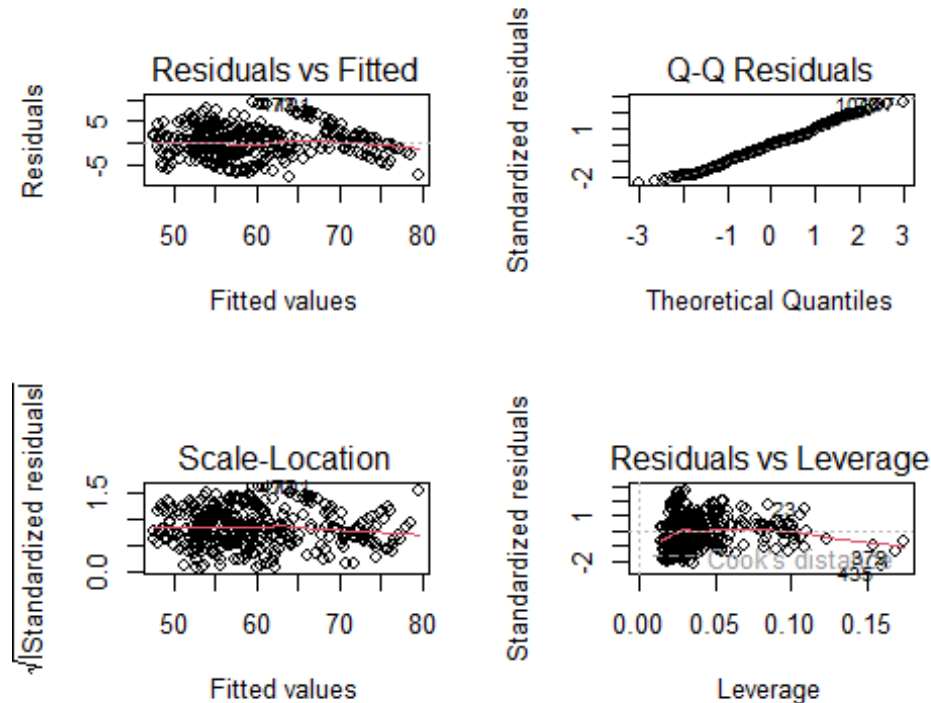
The overall model is significant (p-value < 0.01). The adjusted R-squared is also high (0.8067), implying that the predictors explain 80.67% of the variance in the dependent/outcome variable (Life Expectancy). Year, CO2 emissions and Beer consumption are insignificant predictors of Life Expectancy.

- The intercept is not meaningful in this case.
- On average, life expectancy in countries that aren't least developed is expected to be 6.99 years higher, as compared to life expectancy in countries that are least developed, keeping all other variables constant.
- For every individual increase in population, life expectancy is expected to decrease by (3.582×10^{-8}) , keeping all other variables constant.
- For every unit increase in CO2 emissions, life expectancy is expected to decrease by 0.2153 years, keeping all other variables constant. Even though this isn't significant.
- For every 1% increase in health expenditure as a percentage of GDP, life expectancy is expected to decrease by 0.66 years, keeping all other variables constant. This is very opposite with the obvious expectation that increasing health expenditure would increase life expectancy.
- For every unit increase in electric power consumption, life expectancy is expected to decrease by 0.0027 years, keeping all other variables constant.
- For every square unit increase in forest area, life expectancy is expected to decrease by 0.174 years, keeping all other variables constant. This is also opposite with my expectation that, an increase in forest area would improve life expectancy due to cleaner air, biodiversity, and general health benefits linked to green spaces.

- For every one-dollar increase in GDP per ca-pita, life expectancy is expected to increase by 0.000397 years, keeping all other variables constant.
- For every 1% increase in population with access to the internet, life expectancy is expected to increase by 0.16 years, keeping all other variables constant.
- For every 1% increase in military expenditure as a percentage of GDP, life expectancy is expected to increase by 0.14 years, keeping all other variables constant. This matches the expectation that increasing military expenditure would increase the safety levels of a country, mainly from external attacks like terrorism.
- For every 1% increase in population practicing open defecation, life expectancy is expected to decrease by 0.144 years, keeping all other variables constant.
- For every 1% increase in population with access to basic drinking water services, life expectancy is expected to increase by 0.102 years, keeping all other variables constant.
- For every 1% increase in adult population with obesity, life expectancy is expected to increase by 0.31, keeping all other variables constant. This is also opposite with what I expected.

Linear Model Diagnostics

```
# Plot the model results
par(mfrow = c(2, 2))
plot(results)
```



- There are no patterns in the residuals vs fitted plot, implying that there's a linear relationship between the dependent and the predictor variables.

- The normal Q-Q plot closely resembles a straight line along the diagonal, with little discrepancies on the tails. This implies that the residuals are normally distributed.
- There's no pattern in the Scale-Location plot, implying that there's no heteroscedasticity of the residuals.

Nearly all the assumptions for the Linear model are met.

```
# Cross-validate the linear model to see how it generalizes

# Wrap Learner with feature preprocessing
lm_scaled <- makePreprocWrapperCaret(learner = reg_learner, ppc.scale = TRUE,
                                     ppc.center = TRUE)

# Make resampling description
kFold <- makeResampleDesc(method = "RepCV", folds = 7, reps = 30)

# Cross-validate
lmCV <- resample(lm_scaled, reg_task, resampling = kFold,
                measures = list(rmse, rsq),
                show.info = FALSE)

# View CV results
lmCV

## Resample Result
## Task: training_data
## Learner: regr.lm.preproc
## Aggr perf: rmse.test.rmse=3.6336046,rsq.test.mean=0.7899168
## Runtime: 5.48988
```

An RMSE value of 3.63 is a bit high, the model doesn't perform very well.

Non-linear Regression Models

KNN

```
# Define KNN Learner
kkn <- makeLearner("regr.kknn")

## Loading required package: kknn

##
## Attaching package: 'kknn'

## The following object is masked from 'package:caret':
##
##      contr.dummy

# Wrap Learner with feature preprocessing
kkn_normalized <- makePreprocWrapperCaret(learner = kkn,
```

```

                                ppc.scale = TRUE,
                                ppc.center = TRUE)

# Define hyperparameter space for tuning k
kknParamSpace <- makeParamSet(makeDiscreteParam("k", values = 1:12))
# Specify search strategy
gridSearch <- makeTuneControlGrid()
# Tune the model
tunedK <- tuneParams(kkn_normalized, task = reg_task,
                    resampling = kFold,
                    par.set = kknParamSpace,
                    control = gridSearch,
                    measures = list(rmse, rsq),
                    show.info = FALSE)

# View tuning results
tunedK

## Tune result:
## Op. pars: k=2
## rmse.test.rmse=0.8194258,rsq.test.mean=0.9891001


```

The optimal value of k is 2. RMSE is lower compared to that of the Linear model. R-squared is also closer to 1, which is good. It is however important to note that KNN can overfit at lower values of k.

```

# Extract model information
knnTuningData <- generateHyperParsEffectData(tunedK)
# Visualize the hyperparameter tuning process
plotHyperParsEffect(knnTuningData, x = "k", y = "rmse.test.rmse",
                    plot.type = "line") + theme_bw()

```



RMSE is lowest at $k = 2$.

```
# Set the optimal value of k for the final model
tunedKnn <- setHyperPars(kknn, par.vals = tunedK$x)
# Train the final model
tunedKnnModel <- train(tunedKnn, reg_task)
```

Random Forest

```
# Define Learner
rf_learner <- makeLearner("regr.randomForest")

# Define hyperparameter space for tuning the model
forestParamSpace <- makeParamSet(
  makeIntegerParam("ntree", lower = 100, upper = 100),
  makeIntegerParam("mtry", lower = 4, upper = 25),
  makeIntegerParam("nodesize", lower = 1, upper = 30),
  makeIntegerParam("maxnodes", lower = 5, upper = 25))

# Make resampling description
kFold <- makeResampleDesc(method = "RepCV", folds = 7, reps = 20)

# Specify search strategy
randSearch <- makeTuneControlRandom(maxit = 50)

# Begin parallelization (Parallel processing speeds up the hyperparameter
tuning process)
parallelStartSocket(cpus = detectCores())
```

```

## Starting parallelization in mode=socket with cpus=4.

# Perform hyperparameter tuning with cross-validation
tunedForestPars <- tuneParams(rf_learner, task = reg_task, resampling =
kFold,
                             par.set = forestParamSpace, control =
randSearch,
                             measures = list(rmse, rsq),
                             show.info = FALSE)

## Exporting objects to slaves for mode socket: .mlr.slave.options

## Mapping in parallel: mode = socket; level = mlr.tuneParams; cpus = 4;
elements = 50.

# Stop parallelization
parallelStop()

## Stopped parallelization. All cleaned up.

# View tuning results
tunedForestPars

## Tune result:
## Op. pars: ntree=100; mtry=20; nodesize=21; maxnodes=25
## rmse.test.rmse=1.9145249,rsq.test.mean=0.9417979

```

The RF model outperforms the Linear model but is outperformed by KNN. It has a lower RMSE value of 1.91, even though not very closer to zero.

```

# Set the optimal hyperparameters for the final model
tunedForest <- setHyperPars(rf_learner, par.vals = tunedForestPars$x)

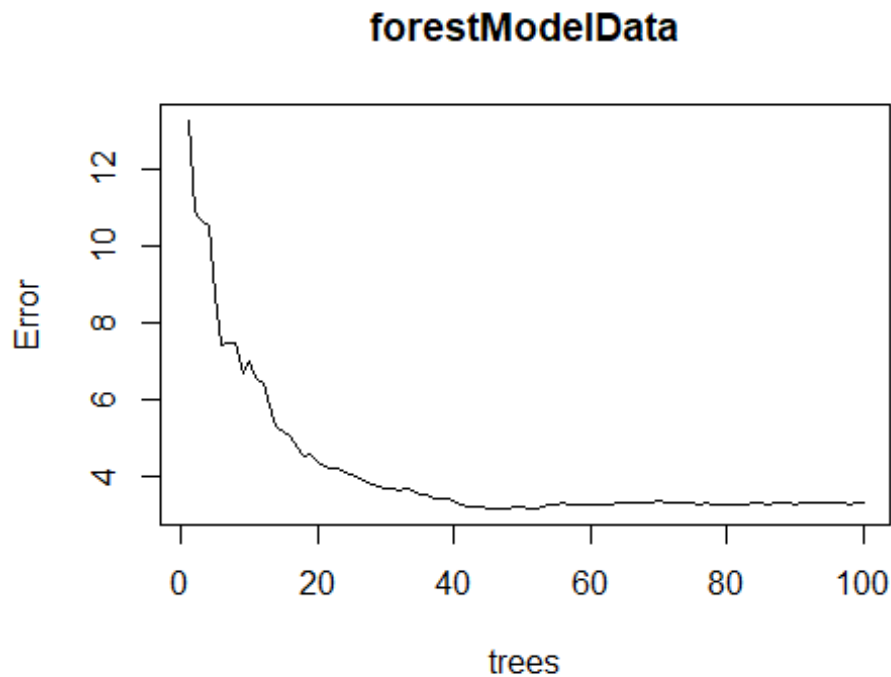
# Train the final model using the optimal hyperparameters
tunedForestModel <- train(tunedForest, reg_task)

## Warning in randomForest.default(x = data[["data"]], y = data[["target"]],
:
## invalid mtry: reset to within valid range

# Extract model information
forestModelData <- getLearnerModel(tunedForestModel)

# Check if there are enough trees in the forest
plot(forestModelData)

```



The out-of-bag error stabilizes after about 75 bagged trees, implying that I have enough trees in the forest.

XGBoost

```
# Define Learner
xgb <- makeLearner("regr.xgboost")

# Define hyperparameter space for tuning
xgbParamSpace <- makeParamSet(
  makeNumericParam("eta", lower = 0, upper = 1),
  makeNumericParam("gamma", lower = 0, upper = 5),
  makeIntegerParam("max_depth", lower = 1, upper = 20),
  makeNumericParam("min_child_weight", lower = 1, upper = 10),
  makeNumericParam("subsample", lower = 0.5, upper = 1),
  makeNumericParam("colsample_bytree", lower = 0.5, upper = 1),
  makeIntegerParam("nrounds", lower = 50, upper = 50))


# Perform hyperparameter tuning with cross-validation
tunedXgbPars <- tuneParams(xgb, task = reg_task, resampling = kFold,
  par.set = xgbParamSpace, control = randSearch,
  measures = list(rmse, rsq),
  show.info = FALSE)

# View tuning results
tunedXgbPars
```

```
## Tune result:  
## Op. pars: eta=0.173; gamma=0.356; max_depth=12; min_child_weight=4.6;  
subsample=0.505; colsample_bytree=0.506; nrounds=50  
## rmse.test.rmse=1.3078593,rsq.test.mean=0.9725991
```

An RMSE value of 1.31 is low and is good (even though not very closer to zero). The XGBoost algorithm outperforms Random Forest and the Linear Regression, but is outperformed by KNN.

```
# Set the optimal hyperparameters for the final model  
tunedXgb <- setHyperPars(xgb, par.vals = tunedXgbPars$x)  
  
# Train the final model using optimal hyperparameters  
tunedXgbModel <- train(tunedXgb, reg_task)  
  
# Extract model information  
xgbModelData <- getLearnerModel(tunedXgbModel)  
  
# Plot the model data to check if there are enough trees in the ensemble  
ggplot(xgbModelData$evaluation_log, aes(iter, train_rmse)) +  
  geom_line() + geom_point() + theme_bw()
```



The curve flattens after the 30th iteration and increasing the number of iterations would not have an effect in the model. This implies that there are enough trees in the ensemble.

Benchmark the KNN, Random Forest and XGBoost model-building processes

```
# Create a tuning wrapper for KNN
kknWrapper <- makeTuneWrapper(kknn_normalized, resampling = kFold,
                             par.set = kknnParamSpace,
                             control = gridSearch,
                             measures = list(rmse, rsq))

# Create a tuning wrapper for RF
forestWrapper <- makeTuneWrapper(rf_learner, resampling = kFold,
                                 par.set = forestParamSpace,
                                 control = randSearch,
                                 measures = list(rmse, rsq))

# Create a tuning wrapper for XGB
xgbWrapper <- makeTuneWrapper(xgb, resampling = kFold,
                              par.set = xgbParamSpace,
                              control = randSearch,
                              measures = list(rmse, rsq))

# Create a list of learners
learners = list(kknWrapper, forestWrapper, xgbWrapper)

# Use holdout cross validation for the benchmarking process
holdout <- makeResampleDesc("Holdout")

# Benchmark
bench <- benchmark(learners, reg_task, holdout, show.info = FALSE)

# View the benchmarking results
bench

##           task.id           learner.id mse.test.mean
## 1 training_data regr.kknn.preproc.tuned      1.056691
## 2 training_data regr.randomForest.tuned      4.790287
## 3 training_data  regr.xgboost.tuned          1.900932
```

According to this benchmarking results, KNN is likely to give me the best-performing model, with a mean prediction error of 1.06

Model Validation

I will use all the four models that I trained to make predictions on test data and assess how they would generalize on new, unseen data. I'll use RMSE as my performance metric.

```
# Make predictions on test data using the Linear model
lmPreds <- predict(Linear_model, newdata = test_data)$data
# Make predictions using KNN model
```



```

knnPreds <- predict(tunedKnnModel, newdata = test_data)$data
# Make predictions using RF model
rfPreds <- predict(tunedForestModel, newdata = test_data)$data
# Make predictions using XGBoost model
xgbPreds <- predict(tunedXgbModel, newdata = test_data)$data

# Calculate test RMSE for each and every model
lm_RMSE <- mean((test_data$life_expectancy - lmPreds$response)^2) |> sqrt()
lm_RMSE

## [1] 3.417389

knn_RMSE <- mean((test_data$life_expectancy - knnPreds$response)^2) |> sqrt()
knn_RMSE

## [1] 0.6378754

rf_RMSE <- mean((test_data$life_expectancy - rfPreds$response)^2) |> sqrt()
rf_RMSE

## [1] 1.873877

xgb_RMSE <- mean((test_data$life_expectancy - xgbPreds$response)^2) |> sqrt()
xgb_RMSE

## [1] 1.194205

```

KNN outperforms all the other algorithms. A test RMSE of 0.638 is low and is good, even though not very closer to zero. On average, the KNN predictions are off by approximately 0.64 years. In other words, I would expect my predictions to be within (plus or minus) 0.64 years of true Life Expectancy.